SMART INDIA HACKATHON 2026

TITLE PAGE

Problem Statement ID: SIH26249

Problem Statement Title: Air Power — Predictive Maintenance & Fleet Availability

Organisation: Ministry of Defence

Idea: FLIGHTLINE — Fleet Availability Twin

Team Name: Code Byte

Code Byte

IDEA TITLE

SMART INDIA HACKATHON 2026

FLIGHTLINE — Fleet Availability Twin

A decision layer that joins the four maintenance silos, predicts how many sortie-days the fleet is about to lose, and names the action that saves the most of them — scored in aircraft-days, not RMSE.

THE PROBLEM :

The statement names fragmentation, not accuracy: *"fragmented and largely reactive … not adequately integrated".*

- Health monitoring, technical records, spares and MRO agencies never meet
- A fault predicted on Tuesday collides with a six-week spares lead time — the aircraft sits on the ground anyway
- RUL models report RMSE; nobody reports aircraft-days

FLIGHTLINE :

ONE GRAPH OF THE FOUR SILOS

tail → LRU → part → stock → depot bay → sorties

PROGNOSIS WITH HONEST INTERVALS

Remaining life as a 90% conformal interval; abstains when unsure

FLEET SIMULATOR + OPTIMISER

90 days, day by day; CP-SAT re-sequences depot inductions and expedites

CANNIBALISATION, PRICED

Every swap returns a net in aircraft-days over 30 days

WHAT THE OFFICER GETS :

- Ranked actions, each scored in aircraft-days gained
- "Why this number?" — the trace of which silo said what
- Availability attribution: why Ao is what it is
- Live what-if: ground aircraft, delay spares, surge sorties
- Printable draft maintenance order

It ranks and prices. It never authorises — the signature stays with the officer.

THE REFRAME: APH

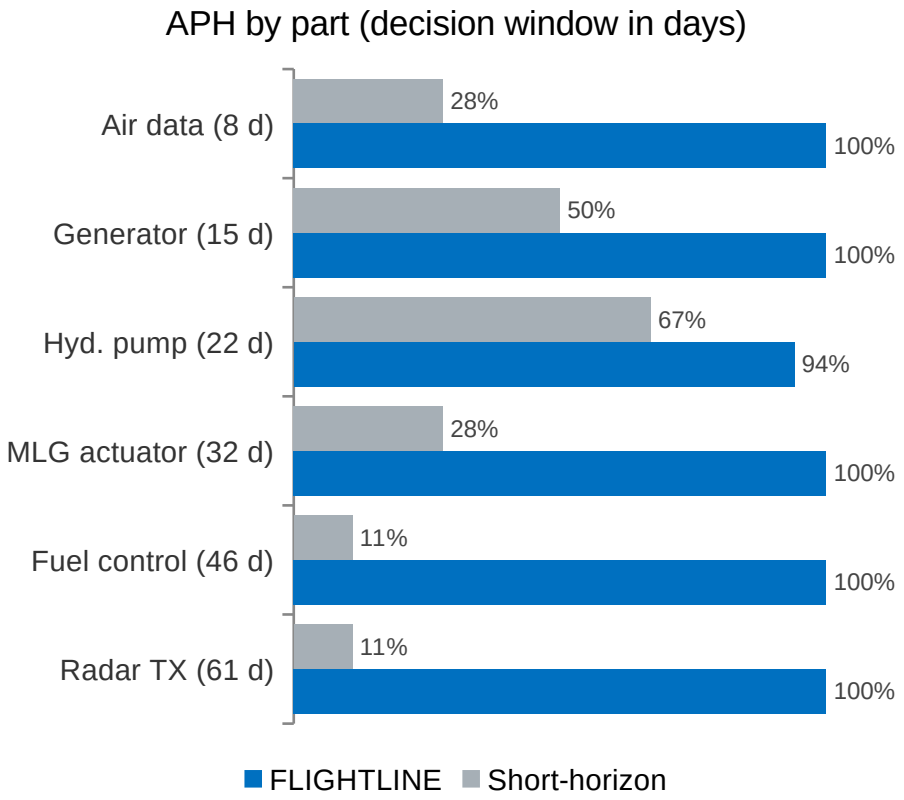
$$APH = P(t_{warning} \leq t_{failure} - (\text{spares lead} + \text{depot queue} + \text{bay time}))$$

A prediction is worth nothing if it arrives after the last day you could have acted on it.

Two real models · same training data · same 108 held-out run-to-failure units

Measure	Short-horizon model (RUL capped at 15, literature convention)	FLIGHTLINE (uncapped + conformal)
RMSE on the model's own target	0.9 cycles	**48.6 cycles**
RMSE in the last 15 cycles before failure	3.7 cycles	**7.1 cycles**
APH — failures warned in time	32%	**99%**
Aircraft-days lost to late warnings / 100 failures	877	**9**
Fleet backtest: sorties lost in 90 days (no prognosis: 18)	12	**8**

More accurate near failure — and still loses. The capped model cannot see past 15 cycles (≈22 days); most spares take longer than that to arrive. FLIGHTLINE warns at a median 1.8× the decision window.



Code Byte

TECHNICAL APPROACH

SMART INDIA HACKATHON 2026

HEALTH MONITORING
144 installed LRUs · multivariate sensor histories, 6 operating regimes (N-CMAPSS shape)

TECHNICAL RECORDS
14 phase inspections due · durations · cannibalisation history

SPARES
18 stock locations · lead times 7–60 days · orders in transit

MRO AGENCIES
2 depots · 3 bays · one induction crew per depot per day

FLEET TWIN one NetworkX graph · 200 nodes · 392 edges · tail → installed LRU → part number → stock location → depot bay → sortie plan

1 PROGNOSIS
- HistGradientBoosting on rolling mean / std / slope per sensor + age
- Mondrian conformalised quantile regression → 90% RUL interval
- Abstains when interval > decision window

2 SIMULATOR
- 90 days, day by day; wear ∝ sorties flown
- Base-stock (S−1, S) spares, FIFO
- Every down-day = one labelled "hold" → exact attribution

3 OPTIMISER
- OR-Tools CP-SAT: induction day + depot, expedites
- Max sortie-days, then aircraft-days
- 2 passes + simulator check; deterministic budget

4 DECISION LAYER
- Cannibalisation priced over 30 days
- Spare unlock ranking (monotone replay)
- Actions in aircraft-days + causal trace

5 DELIVERY
- FastAPI + offline dashboard
- Live what-if, re-solves in ≈1 s
- Printable draft maintenance order

TECHNOLOGY STACK

DATA	ML	OPTIMISATION	GRAPH	SERVING	ASSURANCE
SQLite · seeded generator · N-CMAPSS / C-MAPSS loader	scikit-learn HistGradientBoosting · split-conformal	Google OR-Tools CP-SAT	NetworkX	Python · FastAPI · uvicorn · Chart.js (vendored)	77 pytest tests · same seed → same numbers

CANNIBALISATION, PRICED

What GAO-02-86 found (US Navy + Air Force, FY1996–2000)

≈850,000	cannibalisations reported
≈5.3 million	additional maintenance hours
≥ 2×	"at least twice as long as normal repairs"
111 parts	missing from one aircraft
903–1,756 days	four cannibalised FA-18s not flown
> 25%	readiness fell when the Army stopped (1987 test)

Services' own policy: *cannibalisations "aggravate the very same supply problems that they are meant to overcome."* **Load-bearing and corrosive at once.**

Source: GAO-02-86, Military Aircraft: Services Need Strategies to Reduce Cannibalizations, Nov 2001.

So FLIGHTLINE prices it

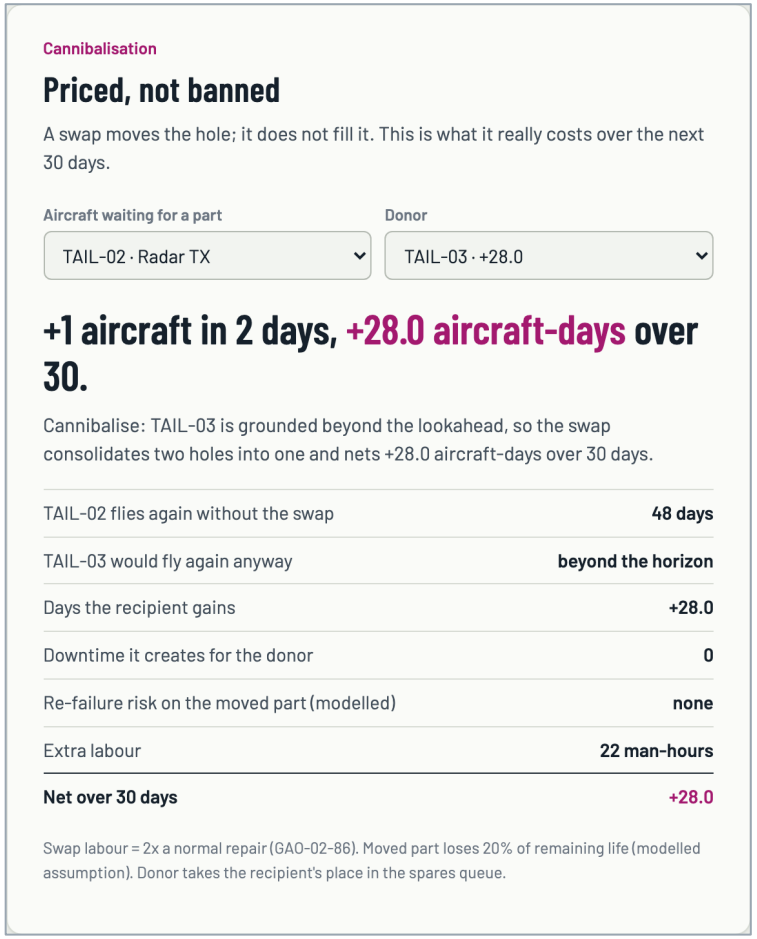
net = recipient gain − donor loss − re-failure

The donor takes the recipient's place in the spares queue: the swap moves the hole, it does not fill it. Swap labour 2× normal (GAO). Moved part loses 20% of remaining life — a modelled assumption, labelled on screen.

Same radar transmitter for grounded TAIL-02:

Donor TAIL-03 grounded for the horizon	**+28.0**
Donor TAIL-06 down 11 more days anyway	**+8.9**
Donor TAIL-03 flying today	**−2.0**

aircraft-days over 30 days



Live card in the prototype (TAIL-01 and TAIL-03 grounded)

Code Byte

THE ABSTENTION GATE

SMART INDIA HACKATHON 2026

We can tell you when we don't know.

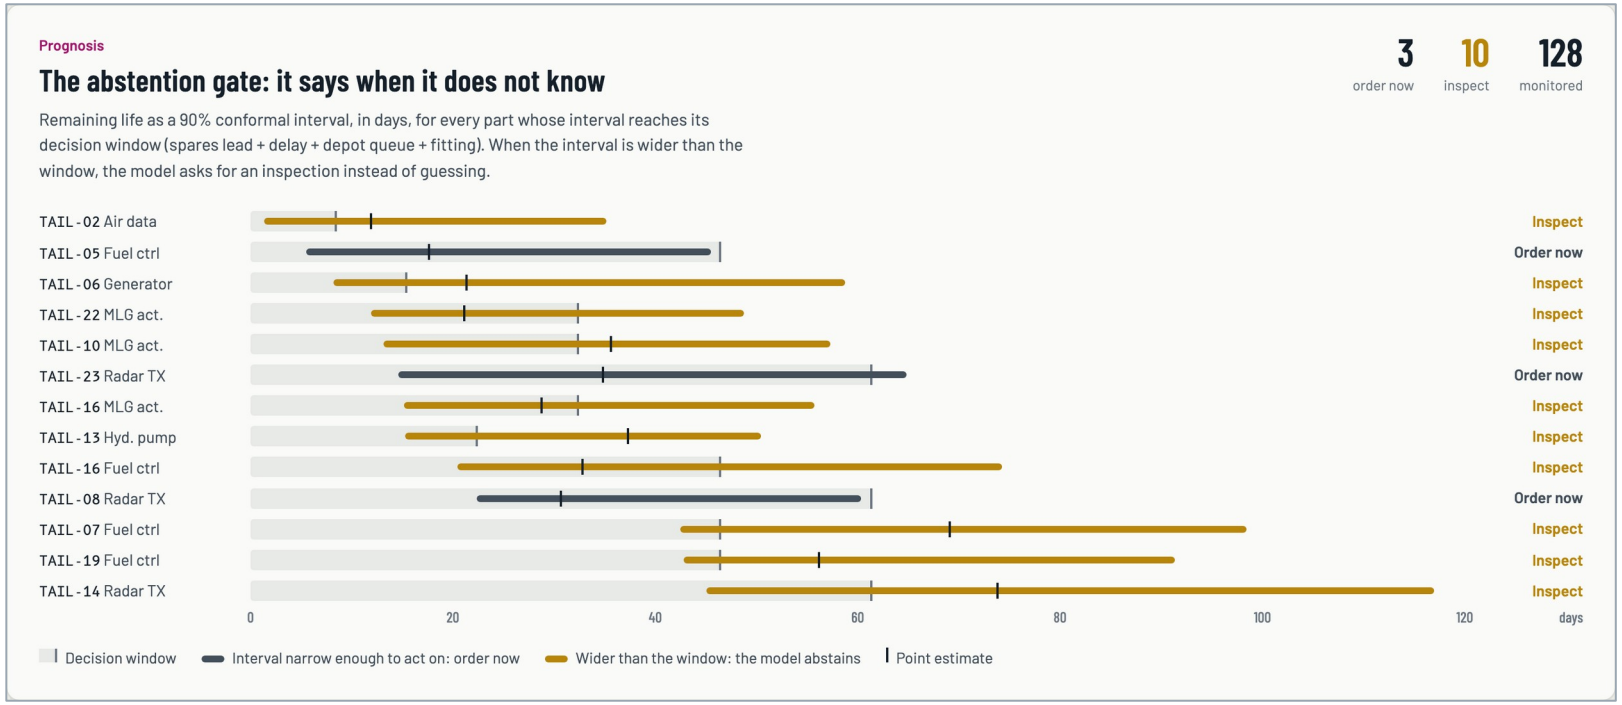
Rule
Abstain when the 90% interval is wider than the decision window **and** its lower edge falls inside it.
→ recommend inspection, not a prediction

91.6% coverage on 1,296 held-out points (target 90%)

94% coverage in the last 20 cycles before failure (85 points)

10 parts flagged for inspection today; 3 confident enough to order; 128 monitored

Mondrian conformalised quantile regression: calibrated separately in three bands of predicted RUL, so the guarantee holds near failure, not just on average.



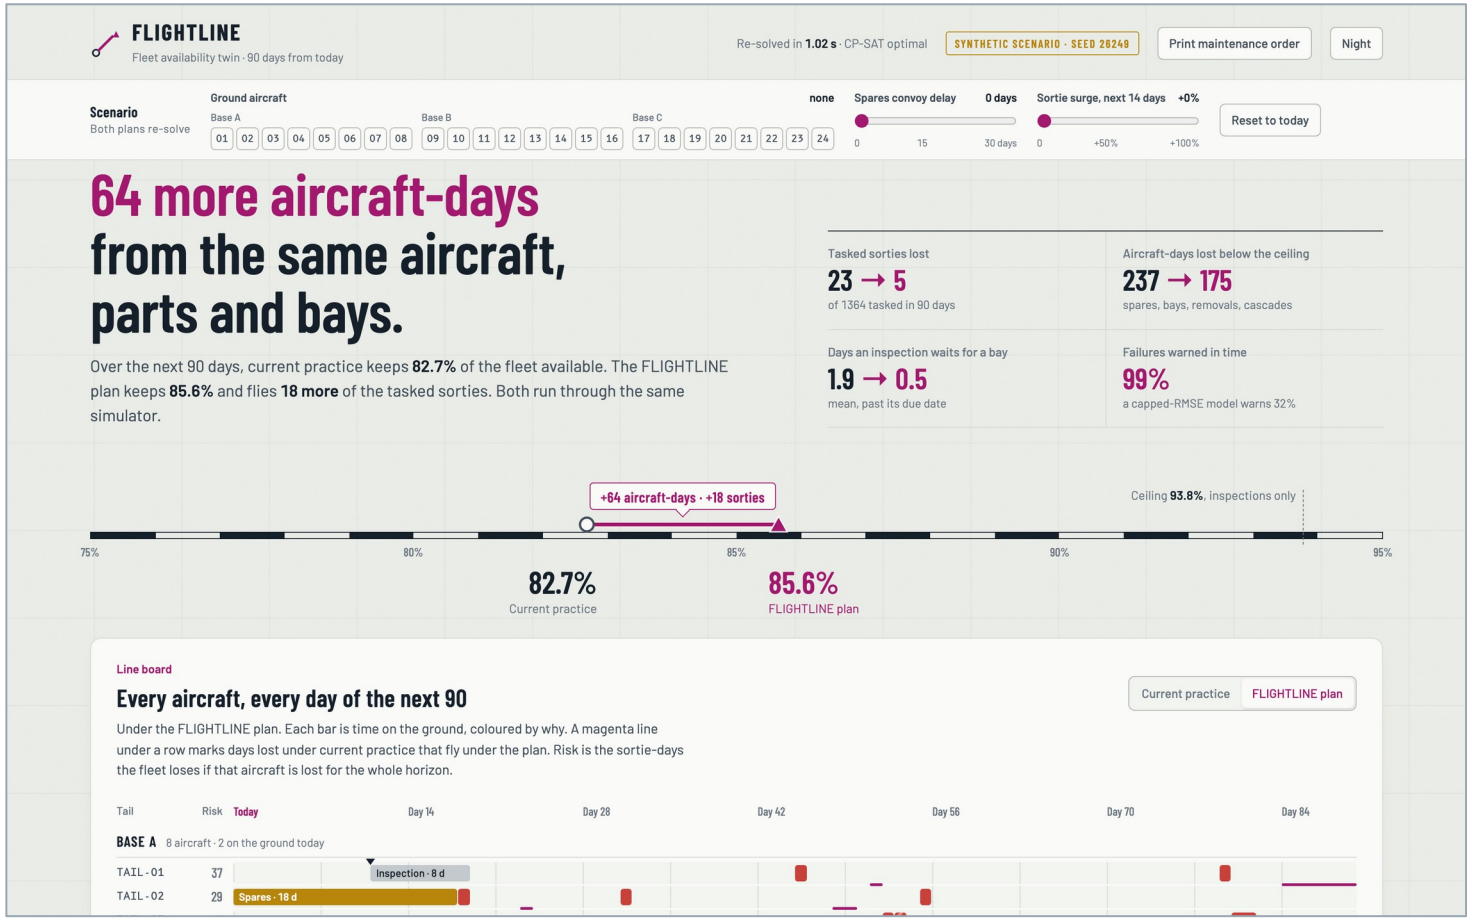
Dark bar = 90% interval narrow enough to act on: order now. **Amber bar** = wider than its decision window (shaded): the model abstains.

Where a false negative can cost aircrew and a false positive wastes a sortie, a system that knows the limits of its own confidence is worth more than one that is slightly more accurate on average.

Code Byte

LIVE DEMO

SMART INDIA HACKATHON 2026



Running prototype · 127.0.0.1:8000 · offline · seed 26249 · re-solves in ≈1.0 s

Ao 82.7% → 85.6% · +64 aircraft-days · +18 sorties flown · sorties lost 23 → 5

Let the panel break it

“Ground two aircraft at Base A.”
Click 01 and 03

Ao falls to 79.0%. Top action becomes: cannibalise the radar from grounded TAIL-03 to TAIL-02, +28.0 aircraft-days.

“Spares convoy is 10 days late.”
Drag convoy delay to 10

Spares share of the gap 8.0 → 9.9 pts; order-now parts 3 → 10; priced cannibalisation +60.

“Surge: 40% more sorties.”
Drag sortie surge to +40%

Sorties lost: 64 under current practice vs 39 under the plan (+83 aircraft-days).

Every action carries “Why this number?” — the trace of which silo said what.

Code Byte

FEASIBILITY AND VIABILITY

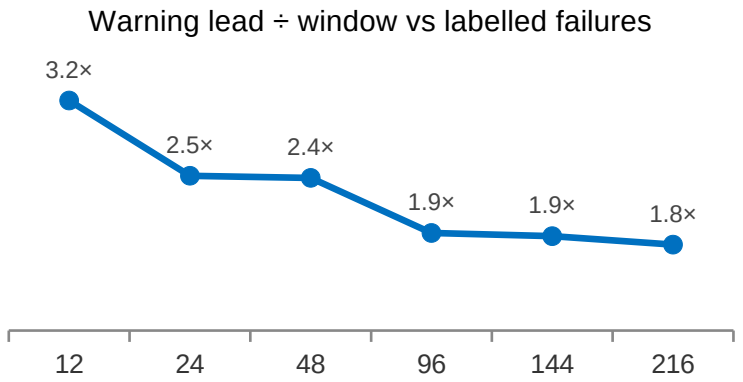
SMART INDIA HACKATHON 2026

HONEST ABOUT DATA

- We have no IAF data and do not pretend to. The badge says SYNTHETIC SCENARIO on every screen
- Degradation: 432 run-to-failure units with the shape of NASA N-CMAPSS
- Drop real N-CMAPSS / C-MAPSS files in data/raw and the model retrains on real engines; logistics stay synthetic and the badge says so
- Decisions never see ground truth; it is used only to backtest

PLATFORM-AGNOSTIC

Swap the degradation model; everything above it is unchanged. Fewer labelled failures do not make it unsafe — the interval widens and it warns earlier:



APH stays 95%–100% from 12 to 216 labelled failures. More data buys efficiency, not safety.

RISKS → MITIGATIONS

Model ≠ real fleet	Decision layer is model-agnostic; conformal recalibrates on a small held-out set
Classified data	Air-gapped on one base server; no cloud, no external API
Trust	Ranks and prices, never authorises; every number traceable
Optimiser approximations	Every plan re-scored by the full simulator

RESEARCH AND REFERENCES

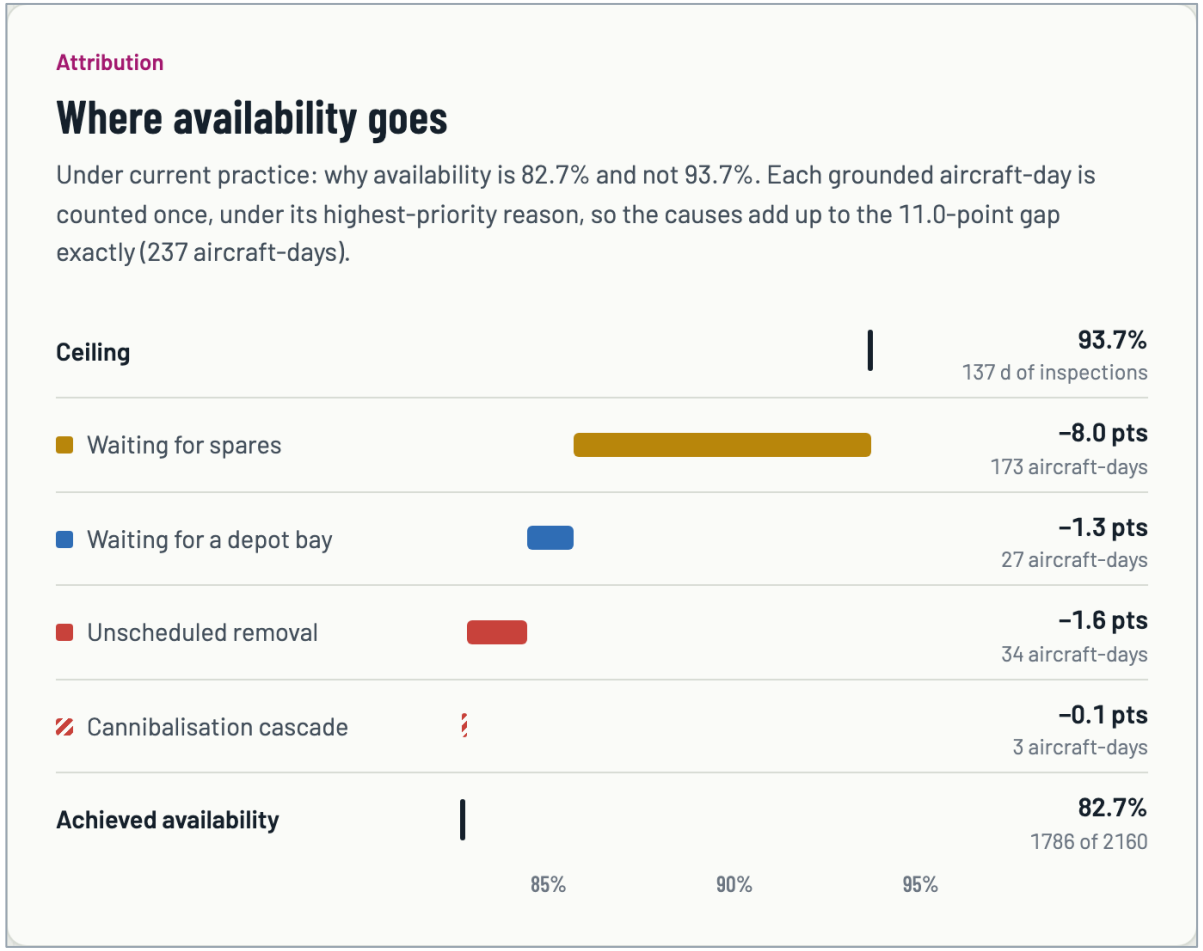
- Arias Chao M. et al., Aircraft engine run-to-failure dataset under real flight conditions (N-CMAPSS), Data 6(1):5, 2021 · Saxena A. et al., C-MAPSS, PHM 2008
- Romano Y., Patterson E., Candès E., Conformalized Quantile Regression, NeurIPS 2019
- US GAO, GAO-02-86, Military Aircraft: Services Need Strategies to Reduce Cannibalizations, 2001 · Sherbrooke C., METRIC, Operations Research 16(1), 1968
- Google OR-Tools CP-SAT solver · scikit-learn HistGradientBoostingRegressor · NetworkX

Code Byte

AVAILABILITY ATTRIBUTION

SMART INDIA HACKATHON 2026

Why is availability 82.7% and not 93.7%?



Cause	Points of Ao	Aircraft-days
Waiting for spares	**8.0**	173
Waiting for a depot bay	**1.3**	27
Unscheduled removal	**1.6**	34
Cannibalisation cascade	**0.1**	3
Total gap below the ceiling	**11.0**	**237**

Every grounded aircraft-day is counted once, under its highest-priority reason, so the causes add up to the gap exactly. It is a unit-tested invariant, not a chart.

This is the screen a Base Maintenance Officer opens at 0700.

Code Byte

IMPACT AND BENEFITS

SMART INDIA HACKATHON 2026

+64

aircraft-days in 90 days on a 24-aircraft fleet (+3.0 pts Ao), same aircraft, parts and bays

23 → 5

tasked sorties lost to unavailability, current practice → plan

1.9 → 0.5

mean days an inspection waits for a depot bay

99%

of failures warned in time to act (capped-RMSE model: 32%)

BASE MAINTENANCE OFFICER

- Why availability is what it is, by cause
- Which depot induction to move, and when

LOGISTICS OFFICER

- Which spare, stocked where, unlocks the most
- Which order to expedite, with its aircraft-days

ENGINEERING OFFICER

- Cannibalisation priced before it is signed
- Inspect-or-order, with the model's own uncertainty

DEPLOYMENT

Air-gapped by construction: one base server, one SQLite file, open-source stack. No cloud call, no LLM, no GPU, no foreign dependency. Nothing leaves the base network.

WHAT WE NEED NEXT

- Labelled removals for one LRU family
- Lead-time history and stock levels by base
- Bay counts and phase intervals for one type; cannibalisation logs